

Assignment 4

Introduction to Parallel Programming

due **Monday 16 October 2023, 23:59**

Instructions

As in the previous assignments, you will need to register in a group on Studium. The group can be different than your previous ones, but, once again, we suggest that you find another student to work together and form a group. In case of trouble, please **contact the assistant** (xiaoyue.chen@it.uu.se).

Submission checklist:

- Submissions must clearly show your name(s).
- Submit a **single** PDF report, as well as all source code files related to the exercises in Studium.
- Solutions must be in C or C++ **with MPI**.
- All source code must compile and run on the **Linux lab machines that have MPI installed**. For the list of the machines you can use for this assignment, see below.
- **Provide instructions for compilation and running, preferably by including Makefile(s)**.
- No source code modifications should be required for reproducing your results.
- Your report must describe relevant details of your solutions and report their performance.

In case you do not reach a working solution, describe the main challenges and proposals to address them. Also, note that this assignment asks you to do benchmarking on the Lab machines and this requires 1) using the machines at a time when they are lightly loaded, and 2) taking multiple measurements to see if there is any variation. Thus, it's not a good idea to leave these tasks for too close to the deadline!

Exercise 0: “Hello, MPI World!” (0 points; no need to submit)

a) Download the provided “Hello!” MPI program from Studium and save it in a file named `hello_mpi.c`. It's time to familiarize yourself on how you can use MPI programs on the lab machines. Some instructions to do this appear below, so follow them step by step.

Start working on the server `gullviva.it.uu.se`. You will be using the servers `tussilago.it.uu.se` and `vitsippa.it.uu.se` as worker nodes.

1. Make sure `ssh` is configured to *not* require passwords. You can check this with the command:

```
ssh your_username@gullviva.it.uu.se
```

which should allow you to login to that machine *without* asking you to type your password.

2. Make sure that you can log into `tussilago` and `vitsippa` from `gullviva` without having to specify a password by following these instructions: http://linuxproblem.org/art_9.html.
3. Create the file “`hosts`”, with the following three lines as content:

```
gullviva.it.uu.se
tussilago.it.uu.se
vitsippa.it.uu.se
```

4. Compile the `hello_mpi.c` program with the following command:

```
mpicc hello_mpi.c -o mpitest
```

5. Run your program using the command

```
mpiexec --hostfile hosts ./mpitest
```

If all goes well, you should be seeing output that looks like:

```
Hello from process 0 of 96
Message from process 1 at gullviva
Message from process 2 at gullviva
...
Message from process 30 at gullviva
Message from process 31 at gullviva
Message from process 32 at tussilago
Message from process 33 at tussilago
...
Message from process 94 at vitsippa
Message from process 95 at vitsippa
```

- b) After you finish with the above task, do something similar with the program that calculates π . (You may need to include some additional header file(s) and comment out some variable(s), e.g., `done`.)
- c) The Web contains a myriad of tutorials for MPI if you need more information to do the exercise that follows. Enjoy reading!

Exercise 1: Sieve of Eratosthenes (2 + 1 + 3 = 6 points in total)

The previous assignments asked you to implement a parallel version of the Sieve of Eratosthenes algorithm for finding prime numbers using Posix threads and OpenMP.

This exercise asks you to use MPI instead of shared-memory parallelism constructs for the parallelization of your solution. Similar constraints as in Assignment 3 apply: If you choose to work in the same (one or two person) group as before, you need to use the program of your previous submission(s) as a basis for this one. If you decide to form a *new* group with another student for this assignment, you can choose *one* of your previous submissions as basis (state which one you used in your report). If you have *not* submitted a program for this exercise before, you can of course write an MPI solution from scratch.

You need to provide and/or measure the following:

1. An MPI program using `MPI_Send()` and `MPI_Recv()` measuring its performance on *one* server.
2. The same program as above measuring its performance on the *three* servers your `hosts` file mentioned above.
3. An MPI program using `MPI_Bcast()` and `MPI_Reduce()` measuring its performance on three servers.

Make sure you use an N which is large enough for the parallelization to make sense.

Besides your code, you need to provide a short section in your report that explains how you modified your solution and reports the speedup curves you got. Was the MPI version easier or more difficult to write? Are the speedups you get the same better or worse (and why)? Refer to previous assignments for more information on how to benchmark your program.

Hint and Warning You can find various tutorials on the web on how to best parallelize the Sieve of Eratosthenes using MPI and also programs that do this. Your solution need not be super sophisticated or optimized, but it should also not be very similar to code “out there”.

Good luck!